

Brian Doran Giffin
Assistant Professor
Structural Engineering & Mechanics
Oklahoma State University

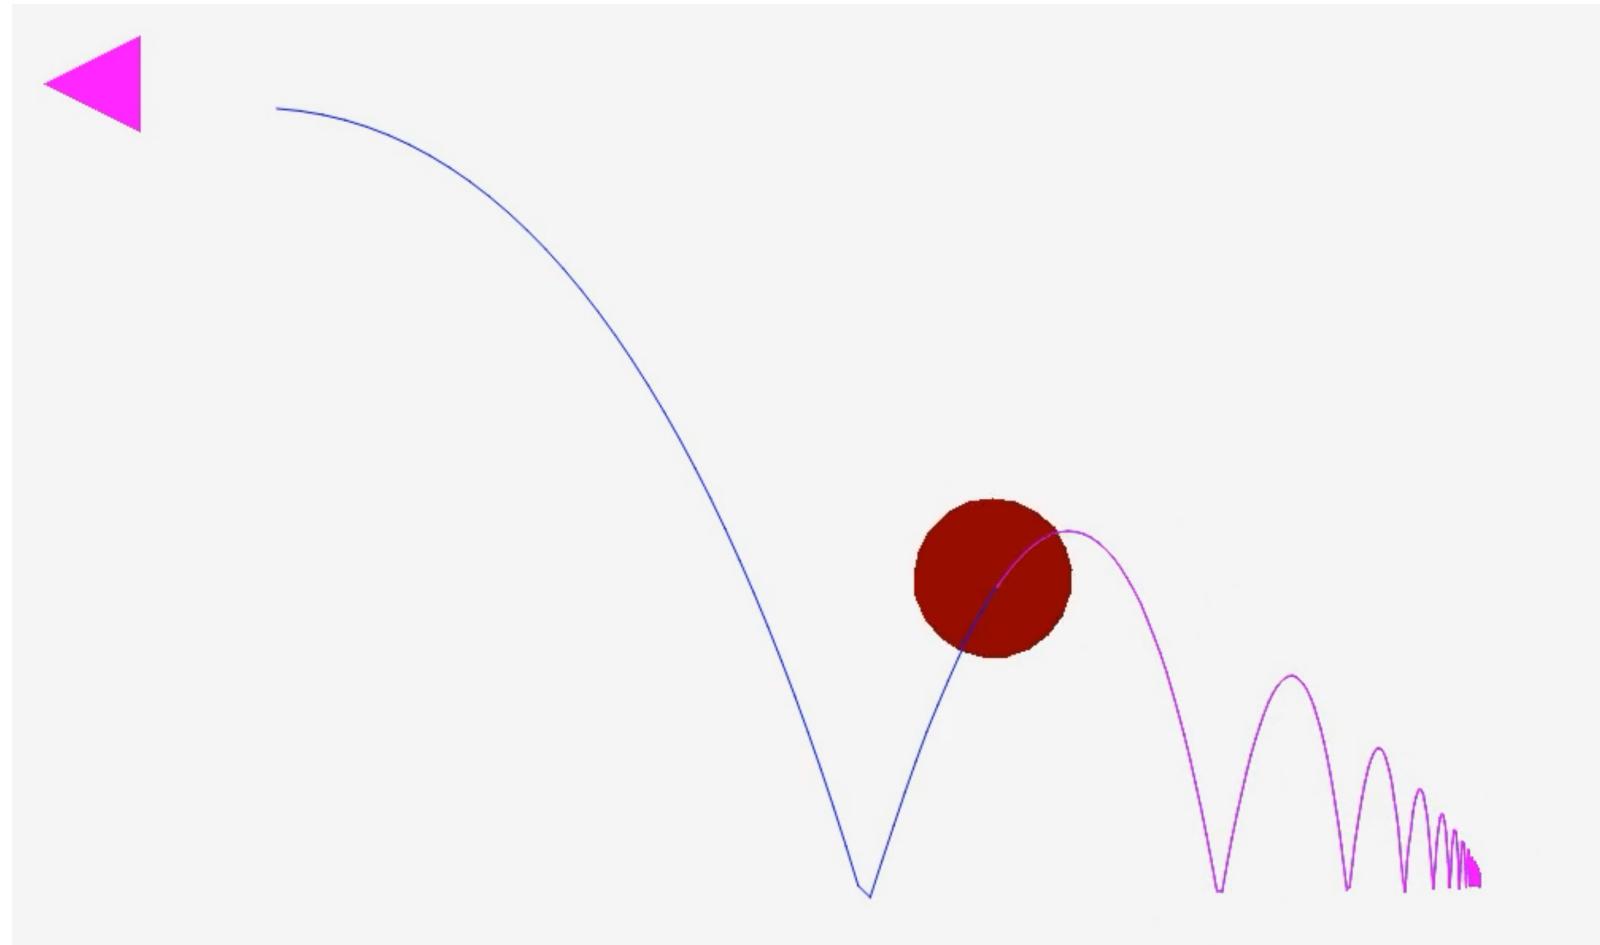


July 20-24, 2025

18th U.S. National Congress
on Computational Mechanics

Computational Methods
for Inverse Problems
and Optimal Experimental Design

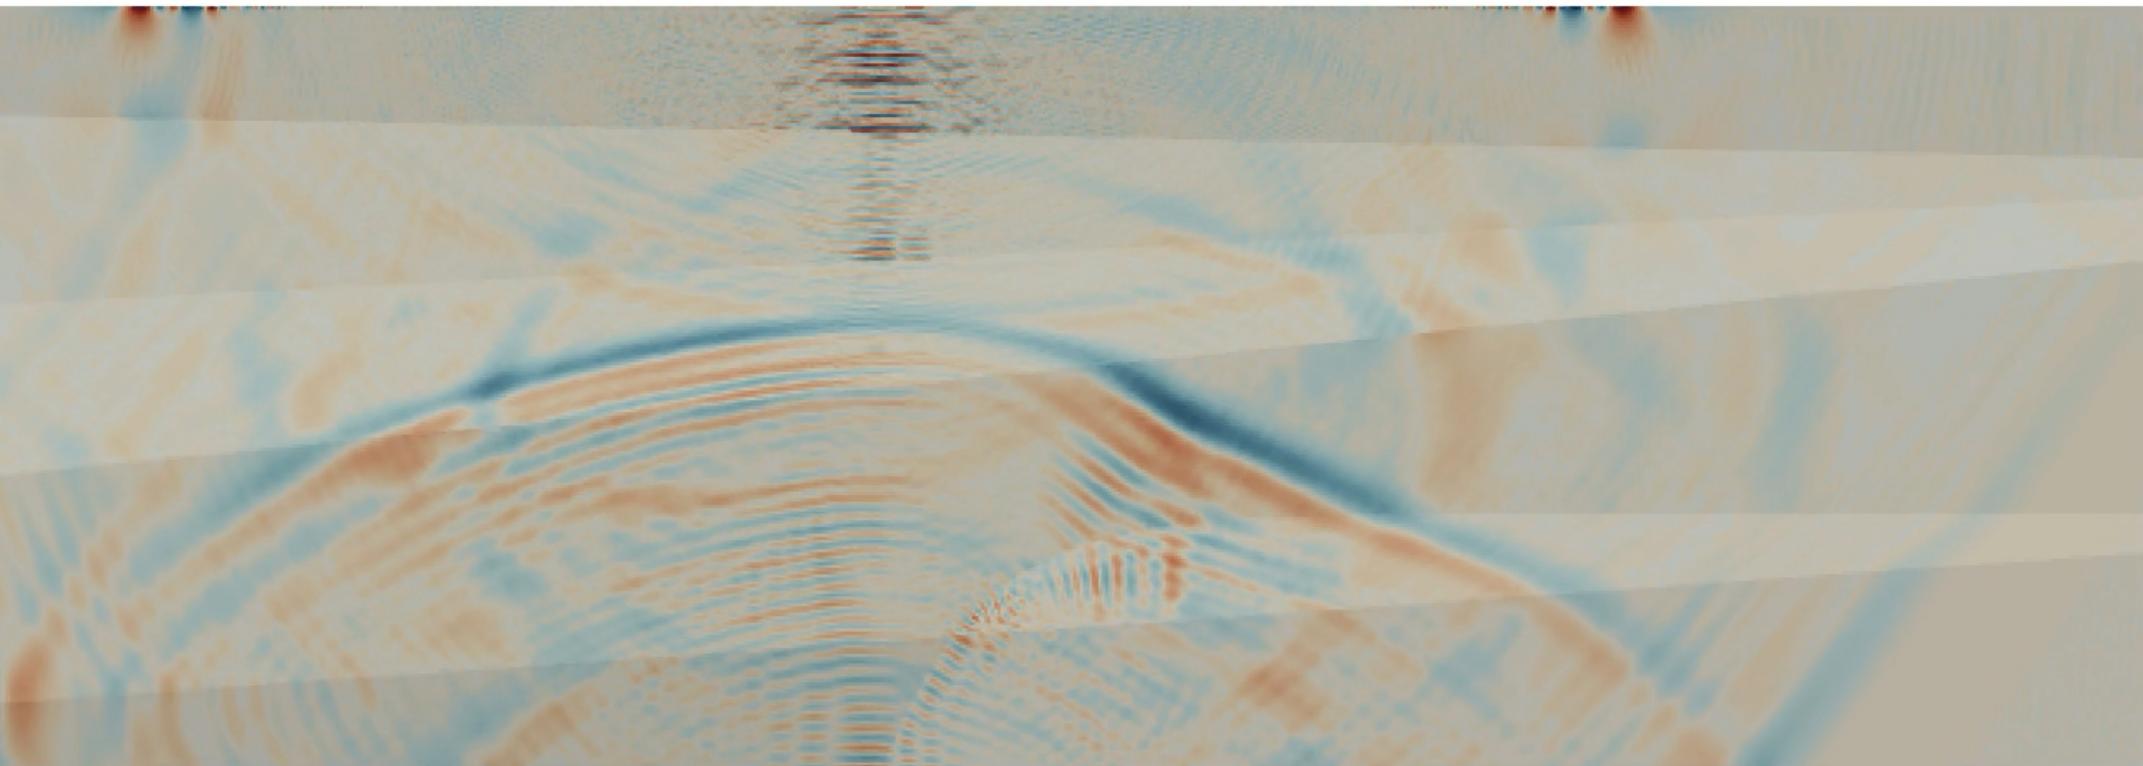
Exactly Bit-Reversible Computational Methods for Memory-Efficient Adjoint Sensitivity Analysis of Dissipative Dynamic Systems



The adjoint state method provides an accurate and efficient means of computing sensitivities for dynamic optimization and inverse problems

Forward problem

t

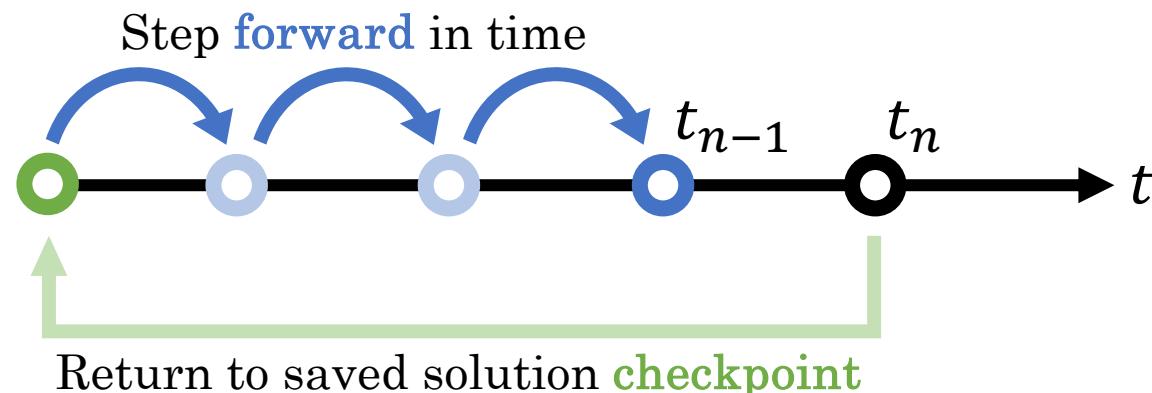


Common implementations of the discrete adjoint method use **checkpointing** to rematerialize forward solution states necessary for backpropagation

Solution of the **adjoint problem** must:

- *Store the forward solution* at all prior time states (**more memory**), or ...
- *Rematerialize forward solution* from “**checkpoints**” (**more computations**)

Forward update:
 $\mathbf{u}_n = \mathbf{f}(\mathbf{u}_{n-1})$

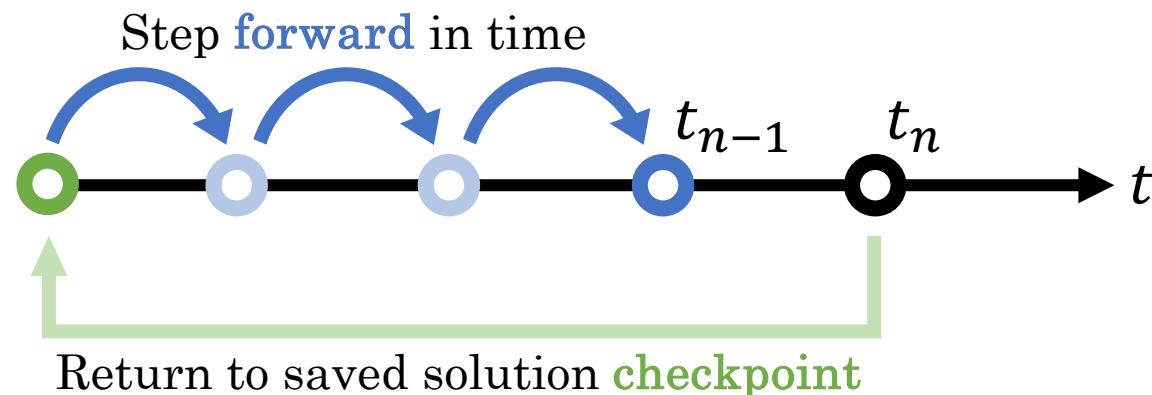


Common implementations of the discrete adjoint method use **checkpointing** to rematerialize forward solution states necessary for backpropagation

Solution of the **adjoint problem** must:

- *Store the forward solution* at all prior time states (**more memory**), or ...
- *Rematerialize forward solution* from “**checkpoints**” (**more computations**)

Forward update:
 $\mathbf{u}_n = \mathbf{f}(\mathbf{u}_{n-1})$



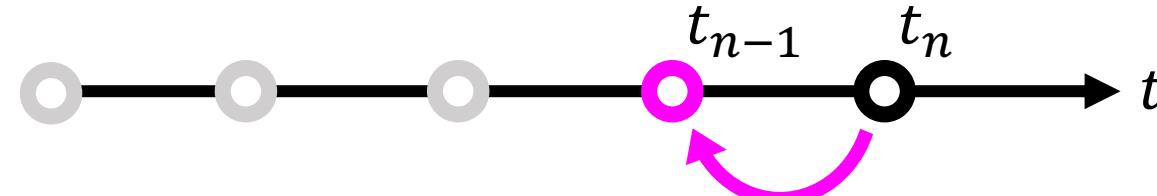
What if we could go *backwards* in time?

Hypothetical solution:

- *Reverse the computations,*
and recover forward solution at each preceding step
- Requires *minimal memory and recomputation*

Backward update:

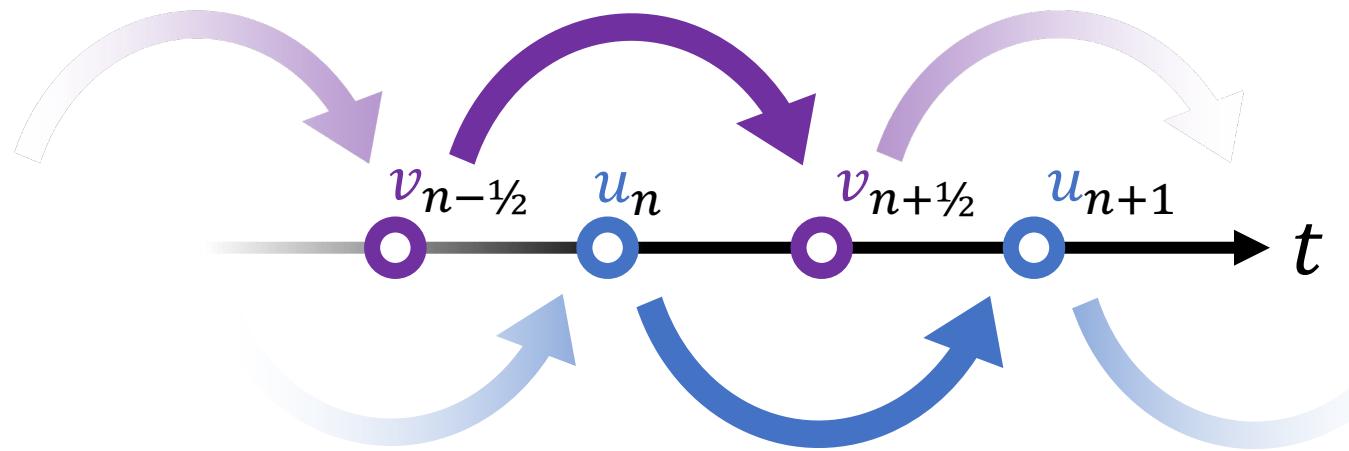
$$u_{n-1} = f^{-1}(u_n)$$



Reverse computations to
step **backwards** in time



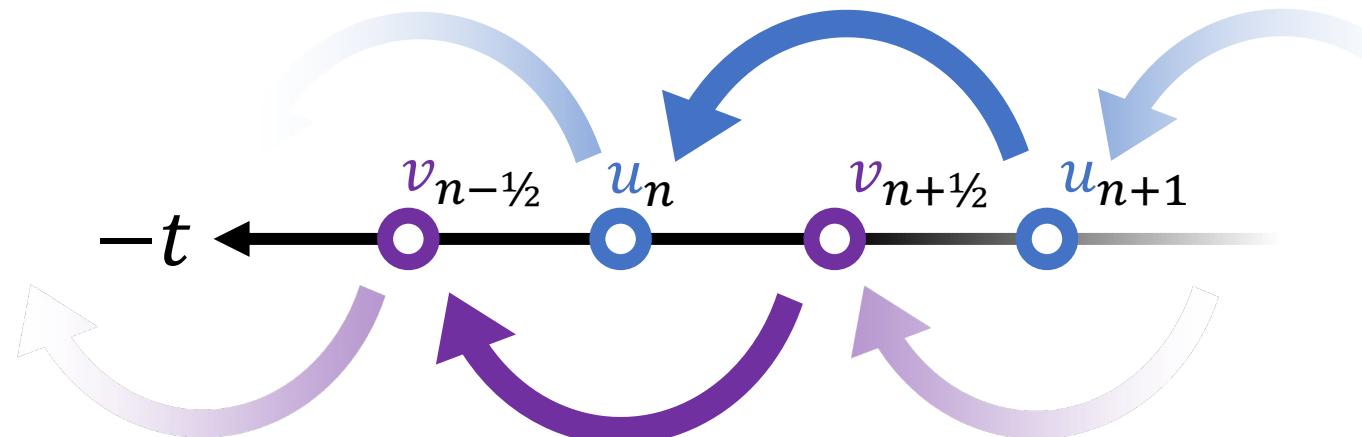
The leapfrog time-integrator is trivially reversible for Hamiltonian systems



Forward update:

$$v_{n+\frac{1}{2}} \leftarrow v_{n-\frac{1}{2}} + a(u_n) \Delta t$$

$$u_{n+1} \leftarrow u_n + v_{n+\frac{1}{2}} \Delta t$$



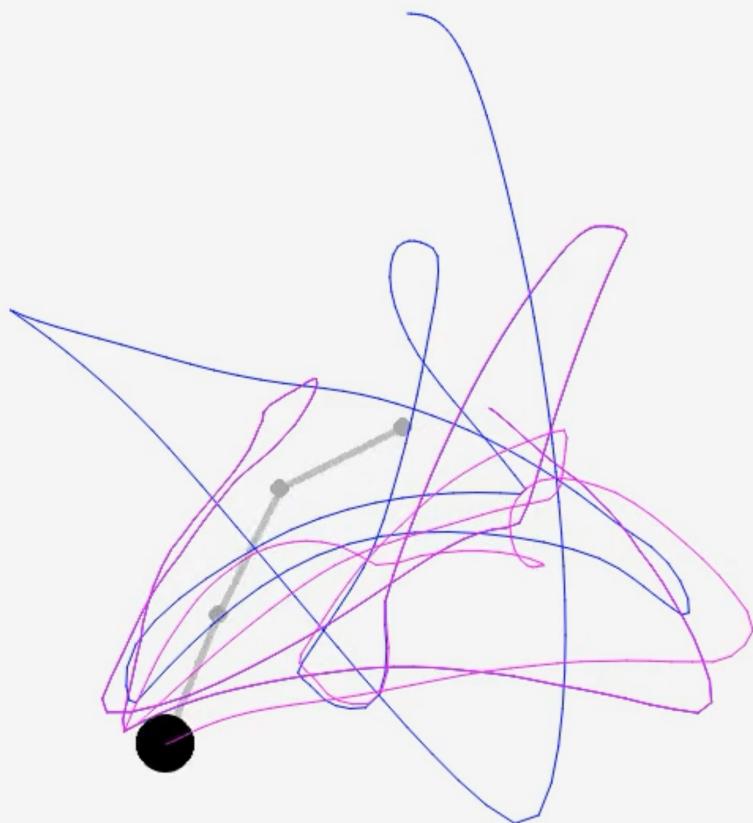
Backward update:

$$u_n \leftarrow u_{n+1} - v_{n+\frac{1}{2}} \Delta t$$

$$v_{n-\frac{1}{2}} \leftarrow v_{n+\frac{1}{2}} - a(u_n) \Delta t$$

Naïve time-reversal leads to inexact rematerialization due to round-off errors from *floating-point arithmetic*

Using *floating-point arithmetic*:



Forward update:

$$v_{n+\frac{1}{2}} \leftarrow v_{n-\frac{1}{2}} + a(u_n) \Delta t$$

$$u_{n+1} \leftarrow u_n + v_{n+\frac{1}{2}} \Delta t$$

Backward update:

$$u_n \leftarrow u_{n+1} - v_{n+\frac{1}{2}} \Delta t$$

$$v_{n-\frac{1}{2}} \leftarrow v_{n+\frac{1}{2}} - a(u_n) \Delta t$$



Represent displacement and velocity degrees of freedom as fixed-width (32- or 64-bit) integers with implied (but differing) problem-dependent radices

$$\begin{array}{l} \text{mantissa} \quad \text{radix} \quad \text{exponent} \\ \downarrow \quad \downarrow \quad \downarrow \\ \text{Fixed } x \in \mathbb{R} \quad x = m \times R^e \quad \text{Integer } m \in \mathbb{Z} \\ \text{e.g. } 7.560239 = 7560239 \times 10^{-6} \end{array}$$

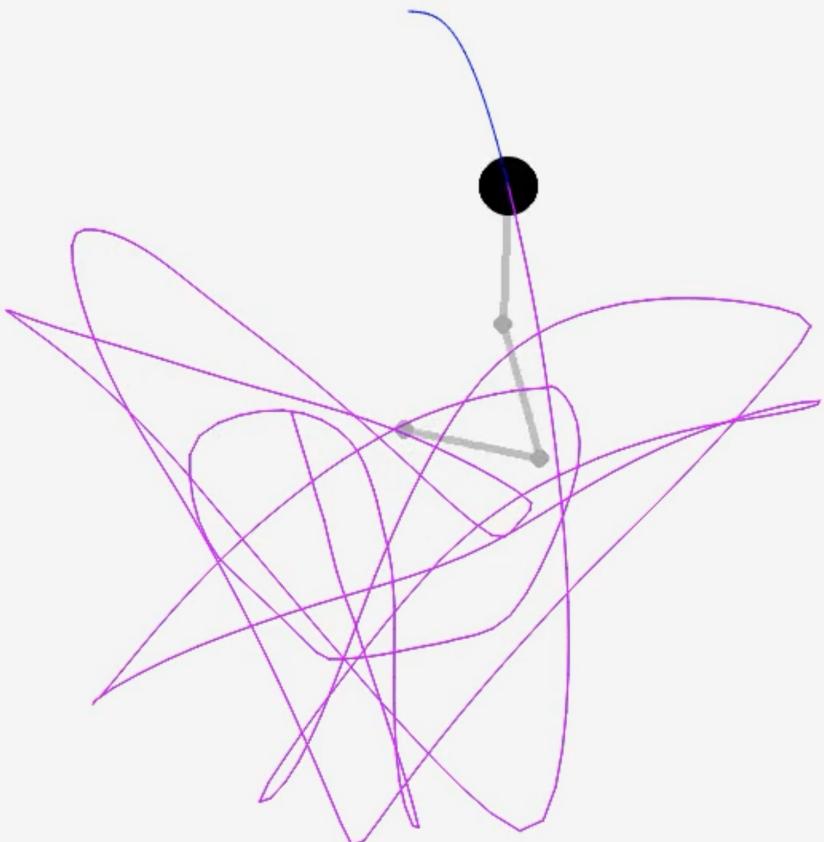
The diagram illustrates the decomposition of a floating-point number x into its components: mantissa, radix, and exponent. The radix is explicitly shown as R^e . The mantissa is highlighted with a bracket labeled "fixed".

$$x \in [-2147.483648, +2147.483647] \quad (32\text{-bit int})$$

$$x \in [-9.223 \dots \times 10^{12}, +9.223 \dots \times 10^{12}] \quad (64\text{-bit int})$$

Round-off errors from addition/subtraction can be eliminated using *fixed-point arithmetic*, resulting in exactly “bit-reversible” time-integration

Using *fixed-point arithmetic*:



Idea previously applied to:

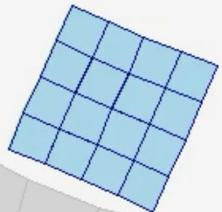
- **Molecular dynamics**
(Levesque and Verlet, 1993)
- **Continuum mechanics**
(Kum and Hoover, 1994)
- **N-body simulations**
(Rein and Tamayo, 2017)
- **Chaotic dynamic systems**
(Jos Stam, 2022)
josstam.com/reversible



For *dissipative* dynamic systems (with damping),
fixed precision arithmetic alone is insufficient to
ensure exact bit-reversibility

Using *fixed-point arithmetic*:

Strain energy: 0.0263
Kinetic energy: 0.0702
Potential energy: 4.5169
Total energy: 4.6134



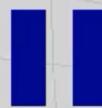
$$M\ddot{u} + C\dot{u} + Ku = F$$

↑
damping

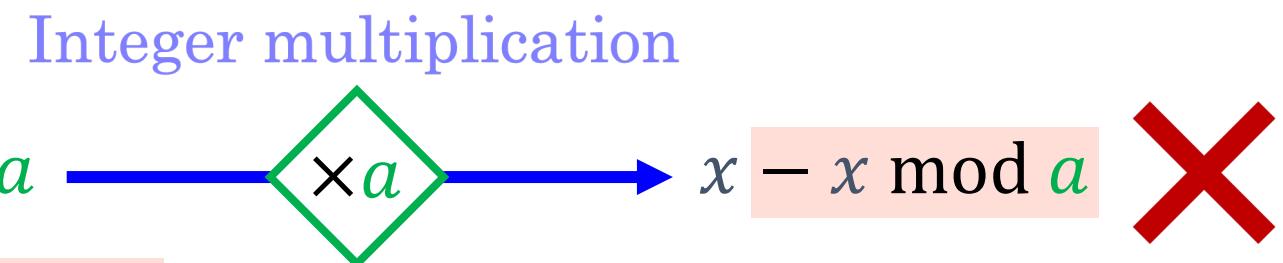
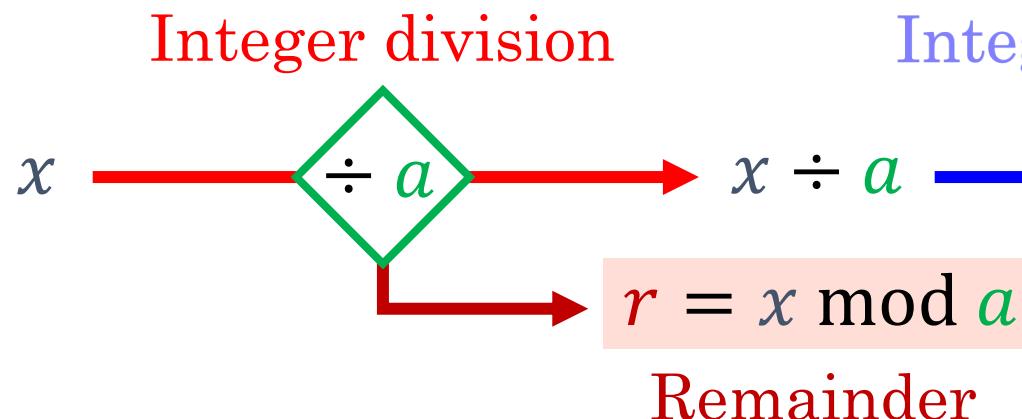
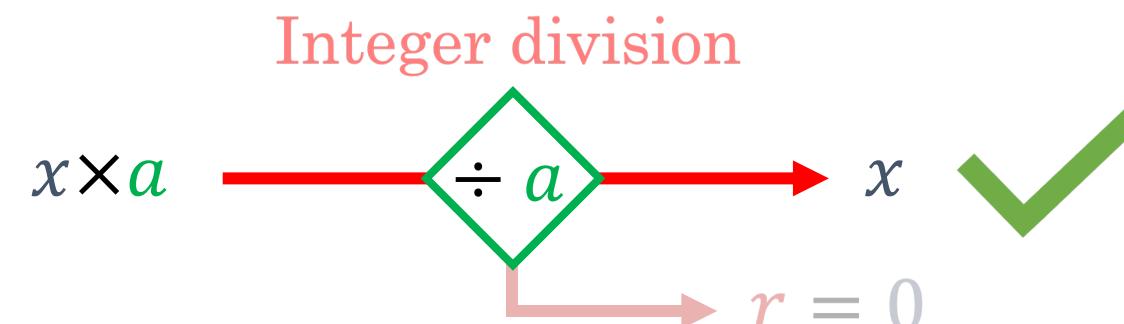
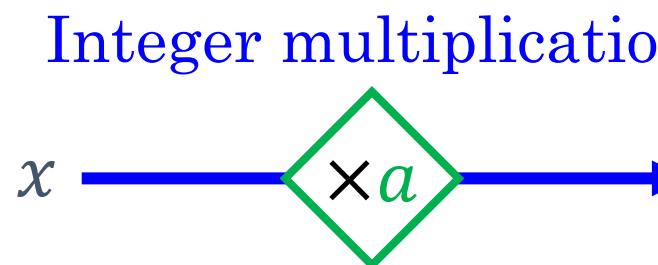
$$\begin{aligned}v_n &\leftarrow v_{n-1/2} + a(u_n)\Delta t/2 \\v_n &\leftarrow v_n \times \frac{M - C\Delta t/2}{M + C\Delta t/2} \\v_{n+1/2} &\leftarrow v_n + a(u_n)\Delta t/2 \\u_{n+1} &\leftarrow u_n + v_{n+1/2}\Delta t\end{aligned}$$

Addition/subtraction: *reversible*

Multiplication/division: *not reversible*

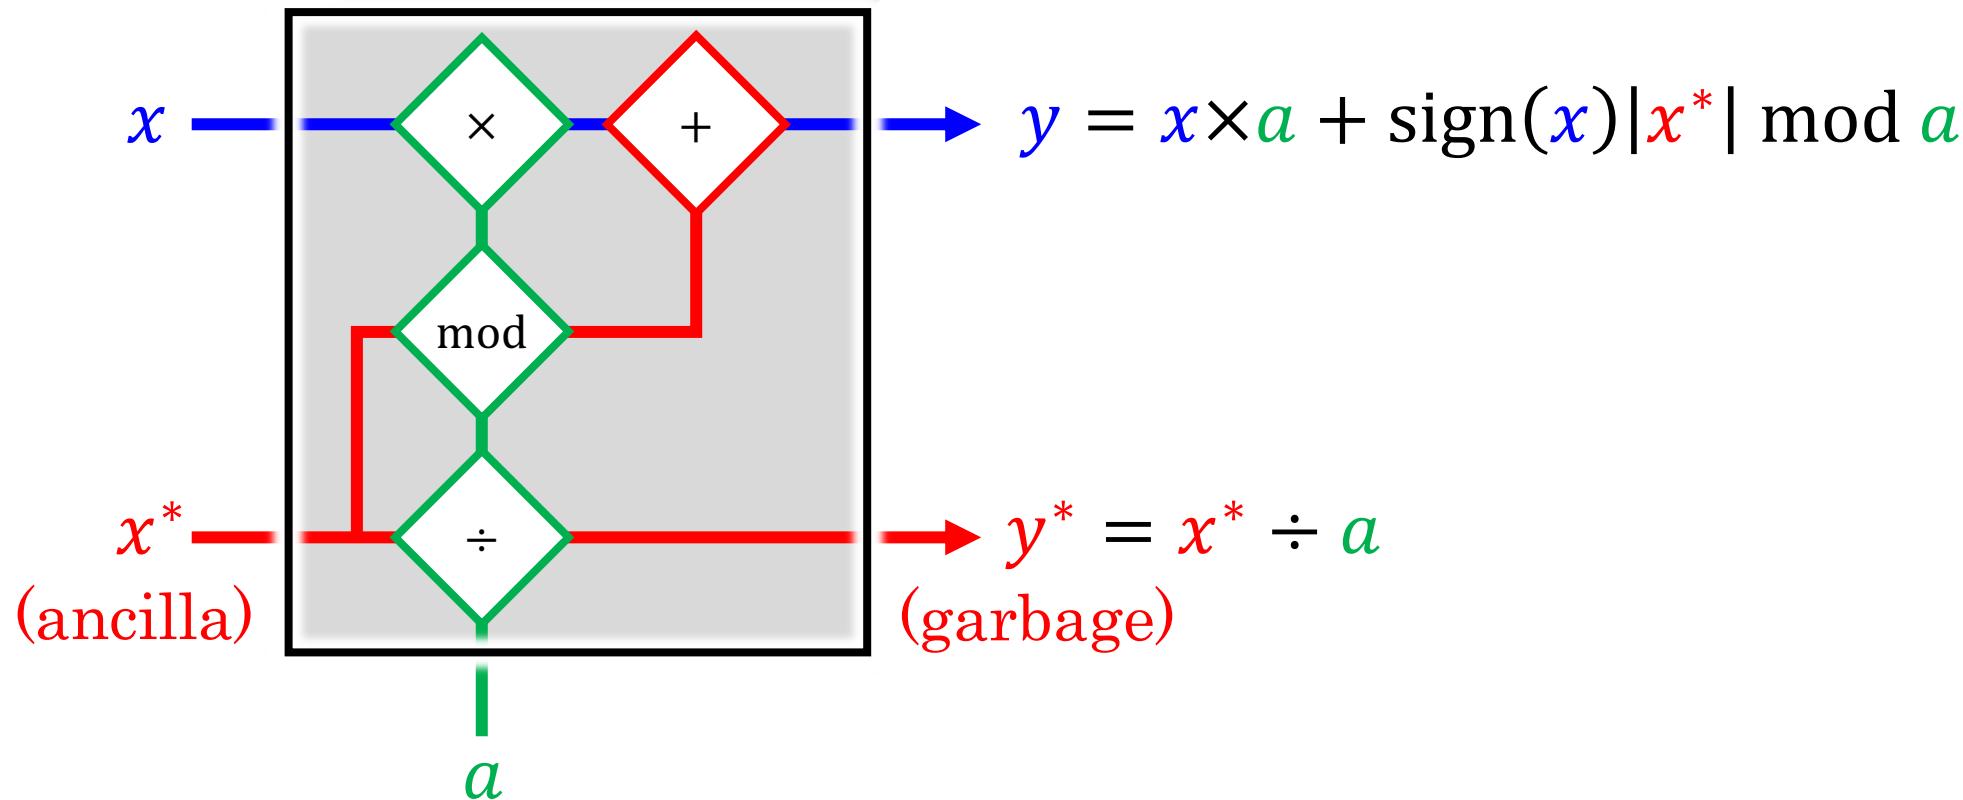


Euclidean division of integers results in permanent loss (“dissipation”) of information in the form of the *remainder* after division



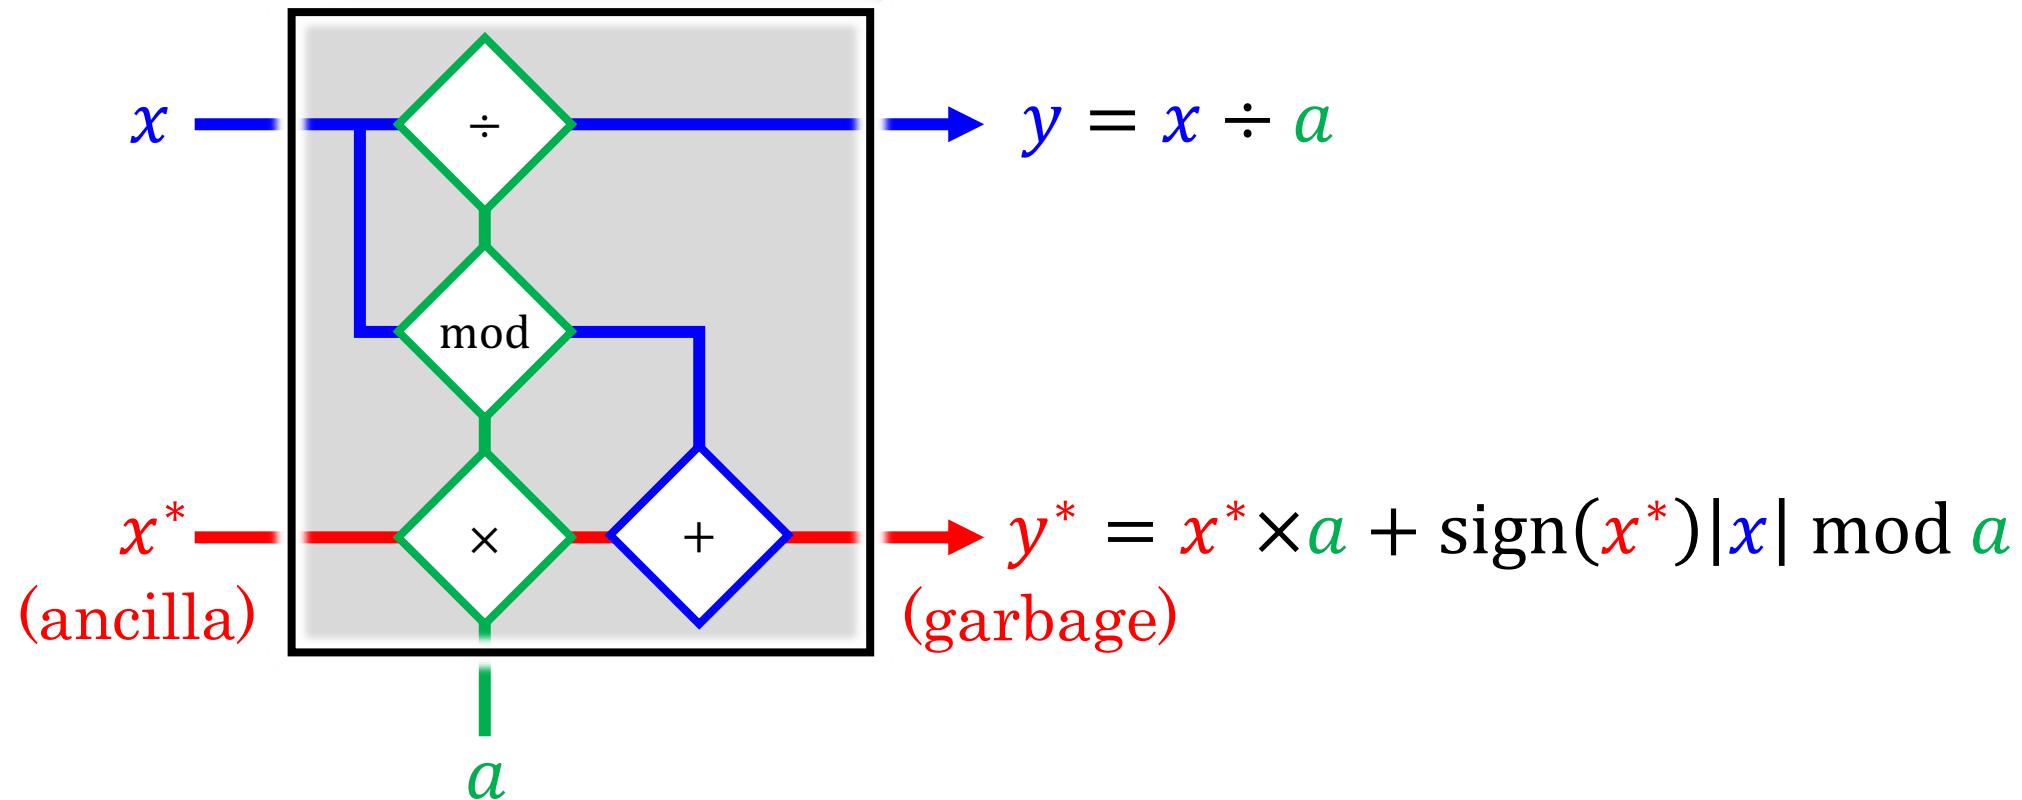
Use the round-off bit buffering approach proposed by Maclaurin et al. (2015) to define a *paired integer multiplication/division operation*

$$(x, x^*)[\times, \div] a = (y, y^*)$$



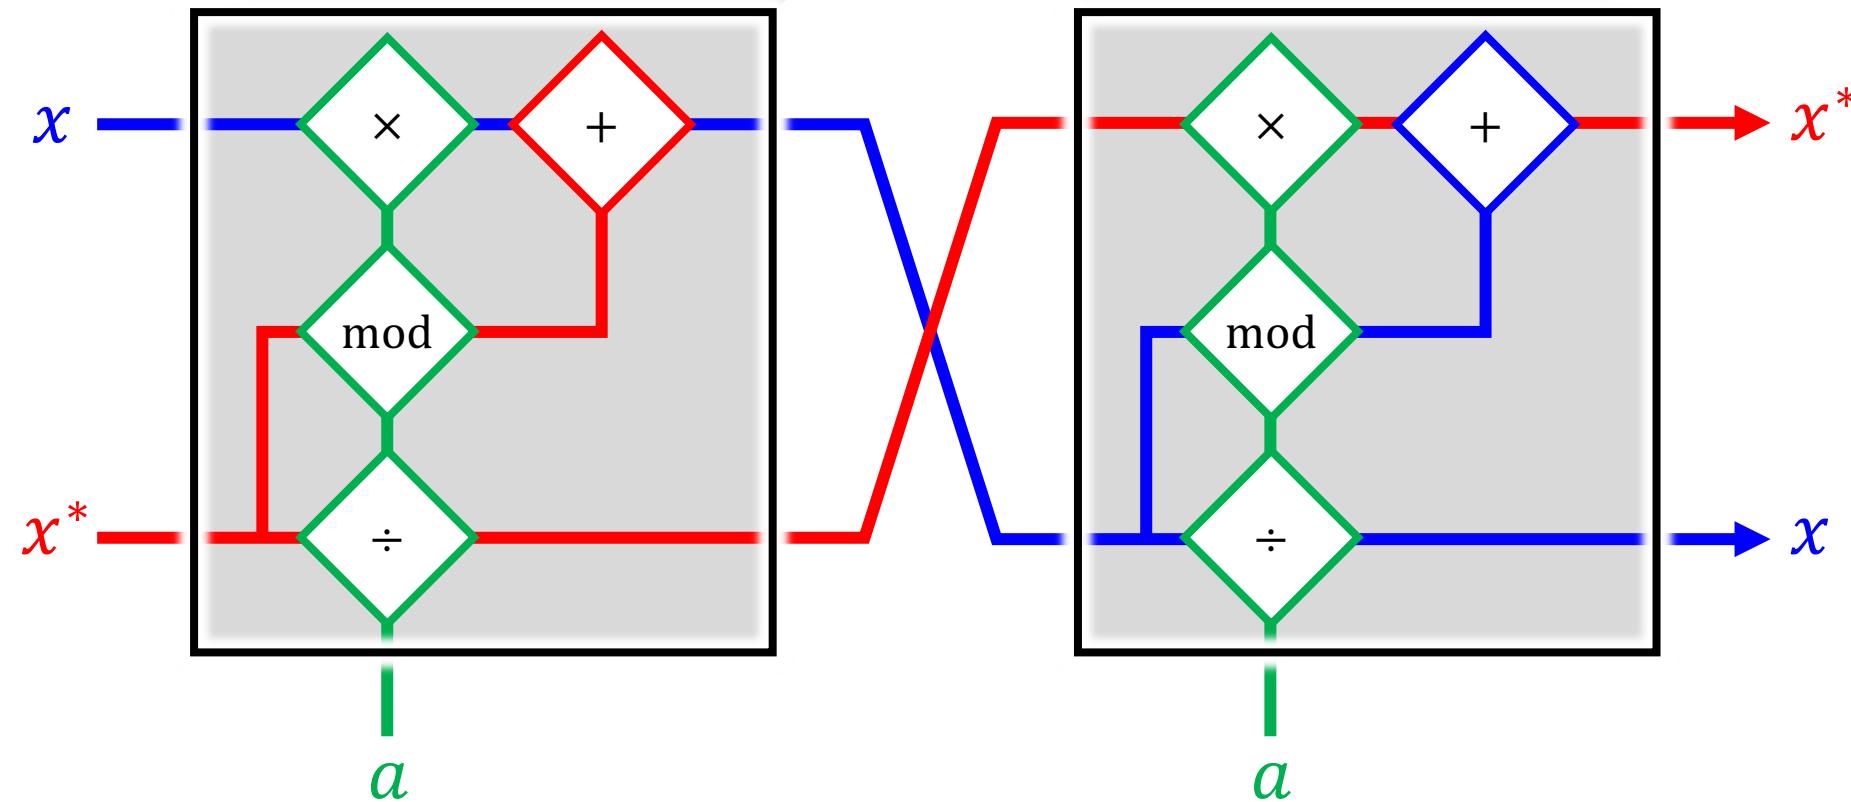
The inverse paired division/multiplication operation is obtained by simply permuting the inputs

$$(x, x^*)[\div, \times]a = (y, y^*)$$



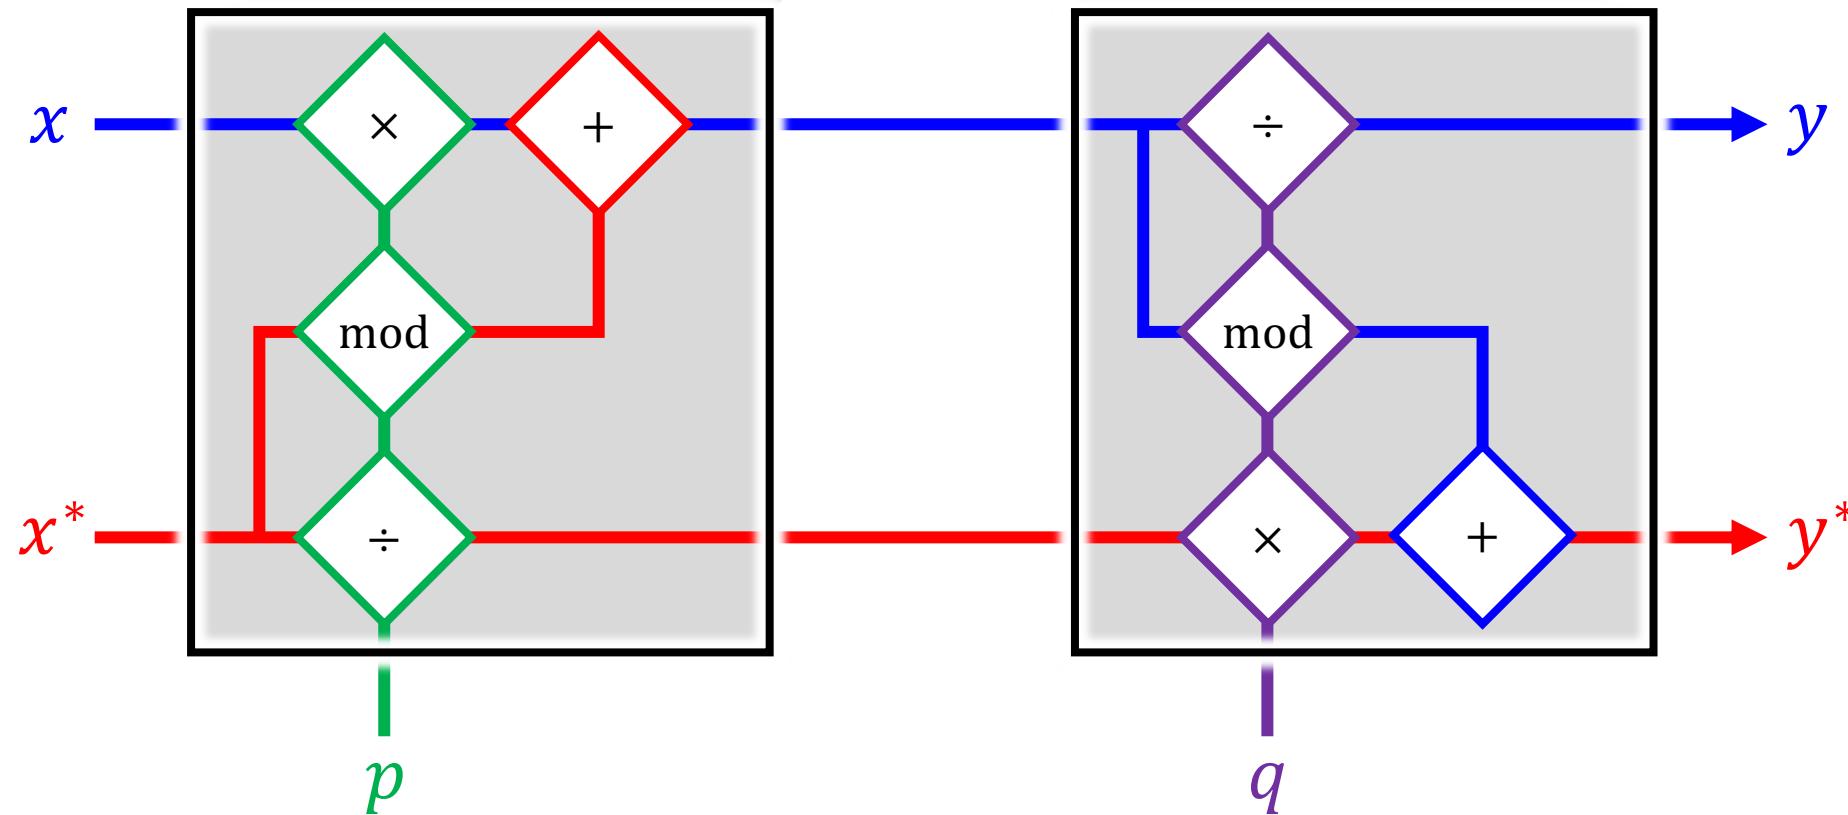
The paired integer multiplication/division operation and its permuted inverse are *exactly bit-reversible*

$$((x, x^*)[\times, \div] a)[\div, \times] a = (x, x^*)$$



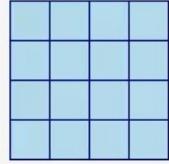
Bit-reversible fixed-point multiplication is carried out by approximating the multiplicand as a *rational number*

$$(x, x^*)[\times, \div] \frac{p}{q} = ((x, x^*)[\times, \div] p)[\div, \times] q = (y, y^*)$$

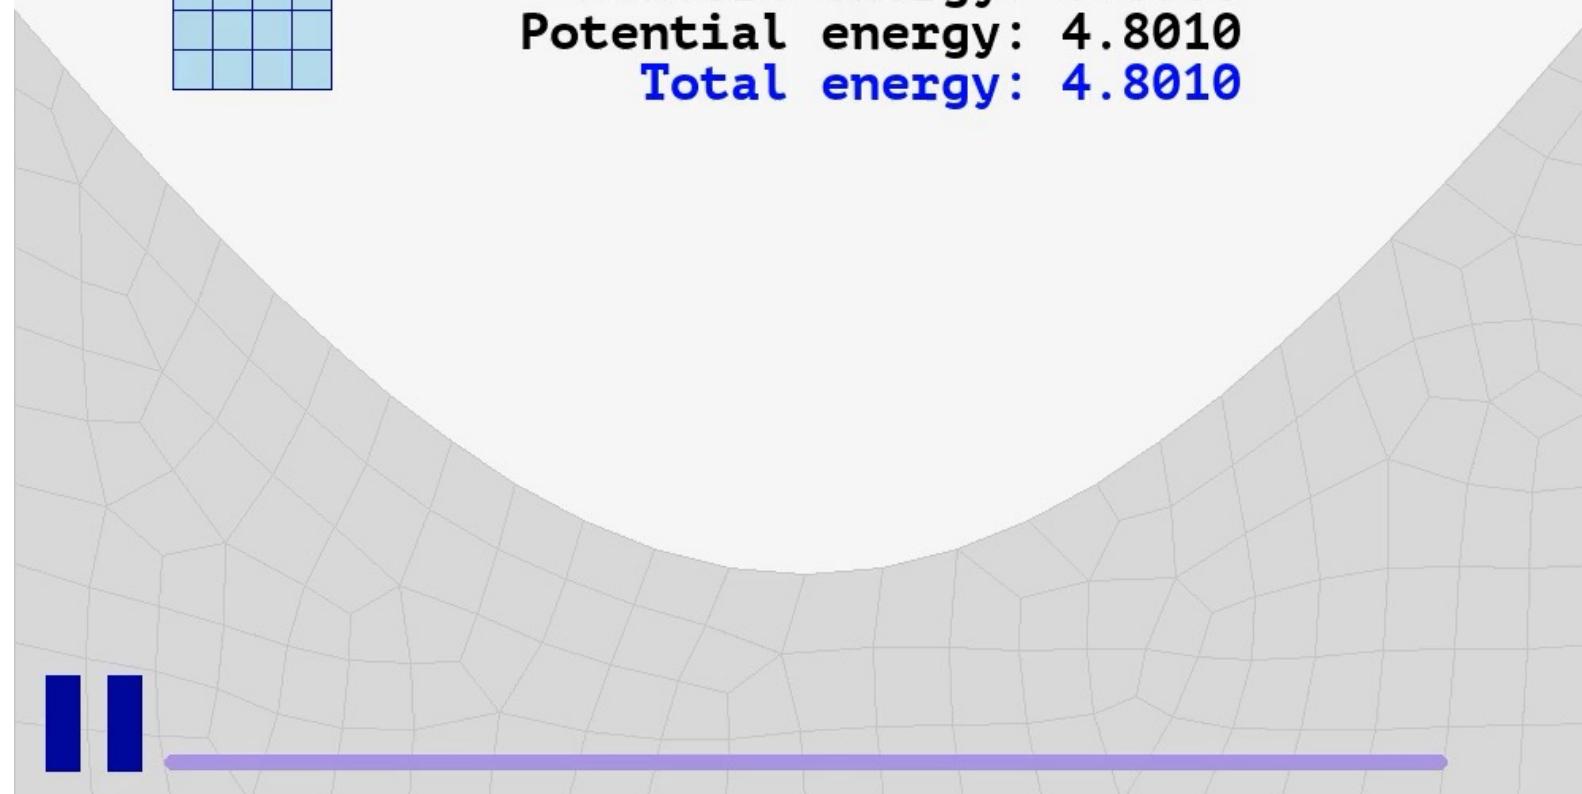


Using paired integer multiplication/division ensures bit-reversibility for dissipative dynamic systems

Using paired integer multiplication/division:



Strain energy: 0.0000
Kinetic energy: 0.0000
Potential energy: 4.8010
Total energy: 4.8010



$$M\ddot{u} + C\dot{u} + Ku = F$$

mass-proportional damping:

$$C = \alpha M$$

$$v_n \leftarrow v_{n-1/2} + a(u_n)\Delta t/2$$

$$(v_n, v_{n+1}^*) \leftarrow (v_n, v_n^*) [\times, \div] \frac{1-\alpha\Delta t/2}{1+\alpha\Delta t/2}$$

$$v_{n+1/2} \leftarrow v_n + a(u_n)\Delta t/2$$

$$u_{n+1} \leftarrow u_n + v_{n+1/2}\Delta t$$

Ancilla velocity state variables: v_n^*

Addition/subtraction: reversible

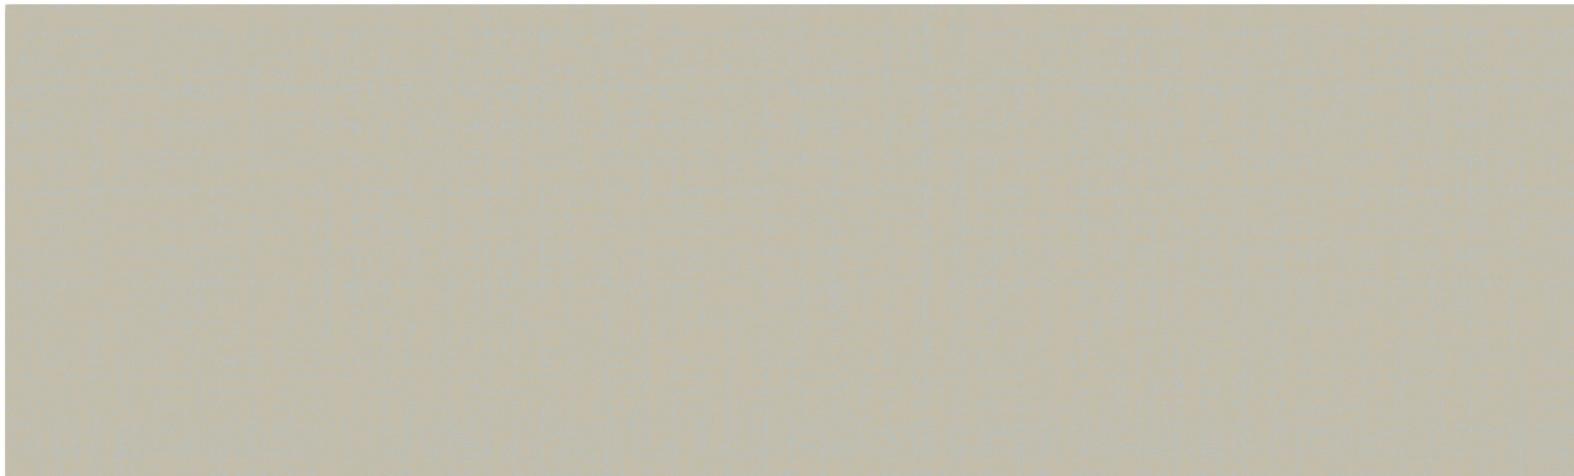
Paired multiplication/division: reversible

Using *paired* integer multiplication/division ensures bit-reversibility for dissipative dynamic systems



Fixed-point arithmetic

unpaired integer
multiplication/division
(not reversible)



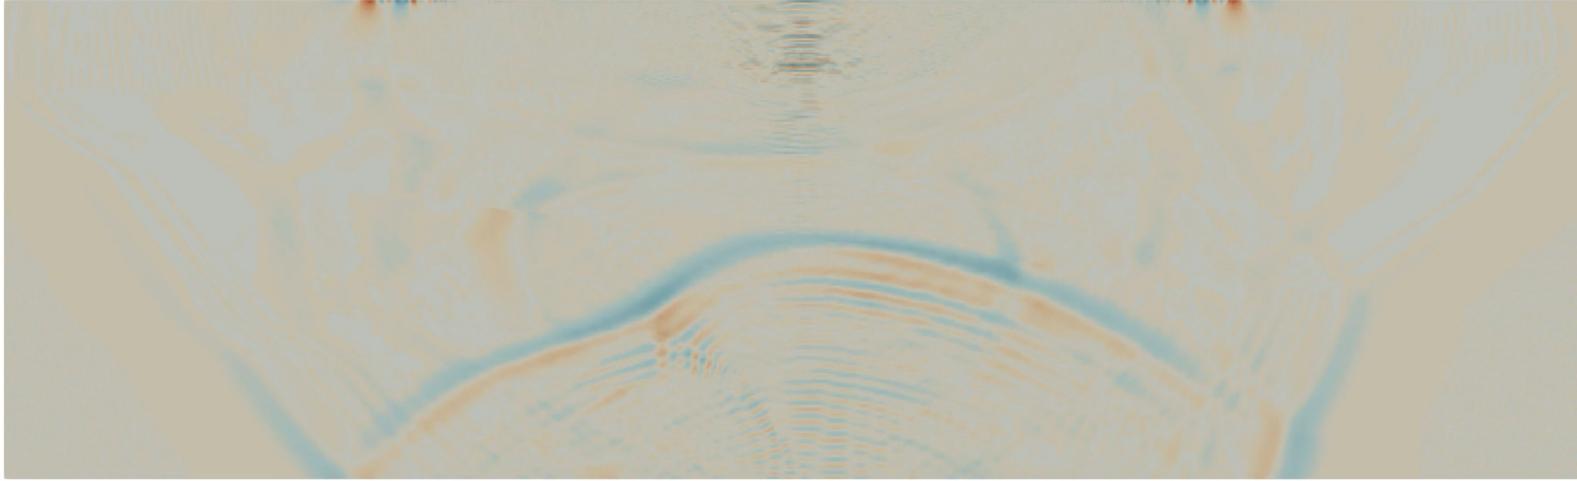
Fixed-point arithmetic

paired integer
multiplication/division
(reversible)

386k degrees of freedom

1000 time steps

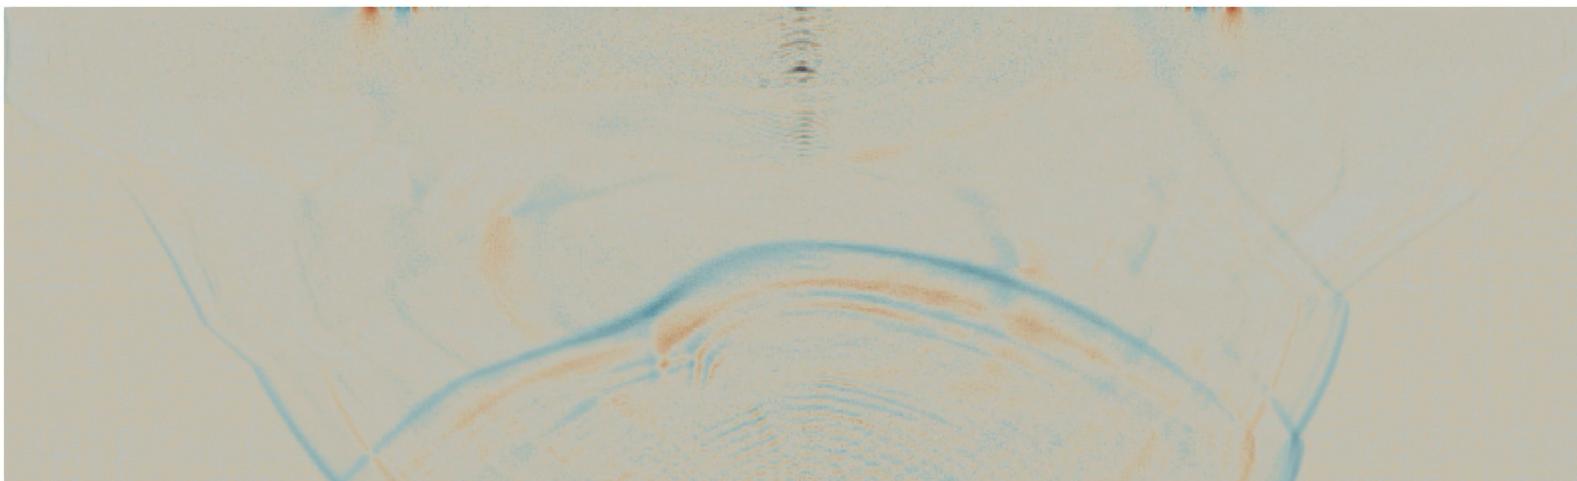
Forward solution accuracy, memory, and run-time performance using paired fixed-point arithmetic is comparable to that of floating-point arithmetic



Floating-point

$1 \times \{64\text{-bit float (double)}\}$

Run-time: 47.3 s



Fixed-point (reversible)

$2 \times \{32\text{-bit int}\}$

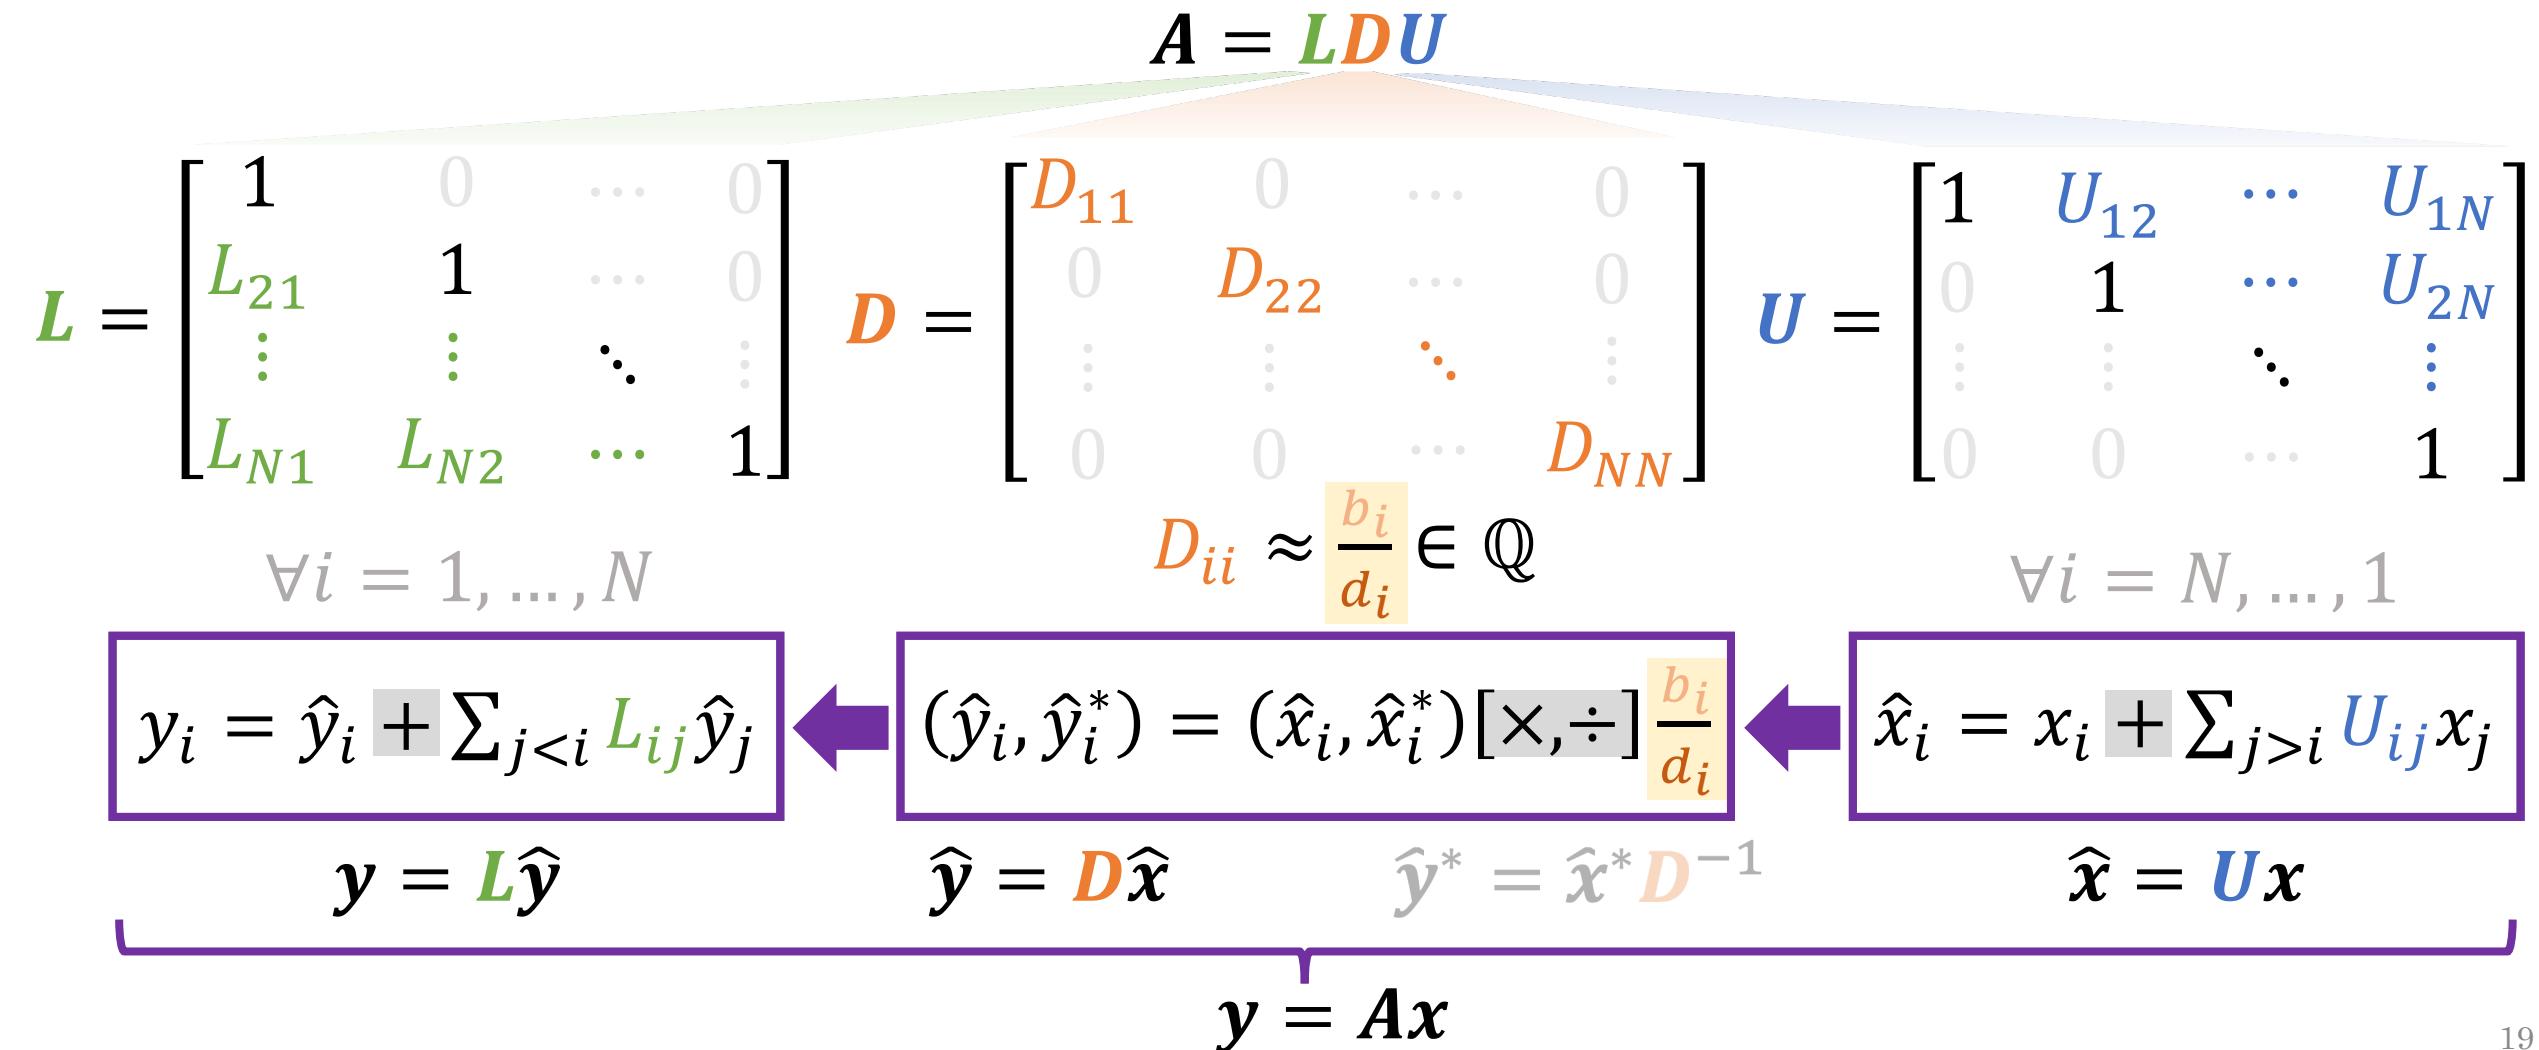
Run-time: 50.2 s

(6% slower)

386k degrees of freedom

1000 time steps

The concept of bit-reversible scalar multiplication/division can be generalized to achieve bit-reversible *matrix multiplication/inversion*



The concept of bit-reversible scalar multiplication/division can be generalized to achieve bit-reversible *matrix multiplication/inversion*

$$A = \textcolor{green}{L} \textcolor{orange}{D} \textcolor{blue}{U}$$

$$\textcolor{green}{L} = \begin{bmatrix} 1 & 0 & \cdots & 0 \\ L_{21} & 1 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ L_{N1} & L_{N2} & \cdots & 1 \end{bmatrix}$$

$$\forall i = 1, \dots, N$$

$$\textcolor{orange}{D} = \begin{bmatrix} D_{11} & 0 & \cdots & 0 \\ 0 & D_{22} & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & D_{NN} \end{bmatrix}$$

$$D_{ii} \approx \frac{b_i}{d_i} \in \mathbb{Q}$$

$$\textcolor{blue}{U} = \begin{bmatrix} 1 & U_{12} & \cdots & U_{1N} \\ 0 & 1 & \cdots & U_{2N} \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & 1 \end{bmatrix}$$

$$\forall i = N, \dots, 1$$

$$\hat{y}_i = y_i - \sum_{j < i} \textcolor{green}{L}_{ij} \hat{y}_j$$

$$\hat{\mathbf{y}} = \textcolor{green}{L}^{-1} \mathbf{y}$$

$$(\hat{x}_i, \hat{x}_i^*) = (\hat{y}_i, \hat{y}_i^*) [\div, \times] \frac{b_i}{d_i}$$

$$\hat{\mathbf{x}} = \textcolor{orange}{D}^{-1} \hat{\mathbf{y}}$$

$$\hat{\mathbf{x}}^* = \hat{\mathbf{y}}^* \textcolor{orange}{D}$$

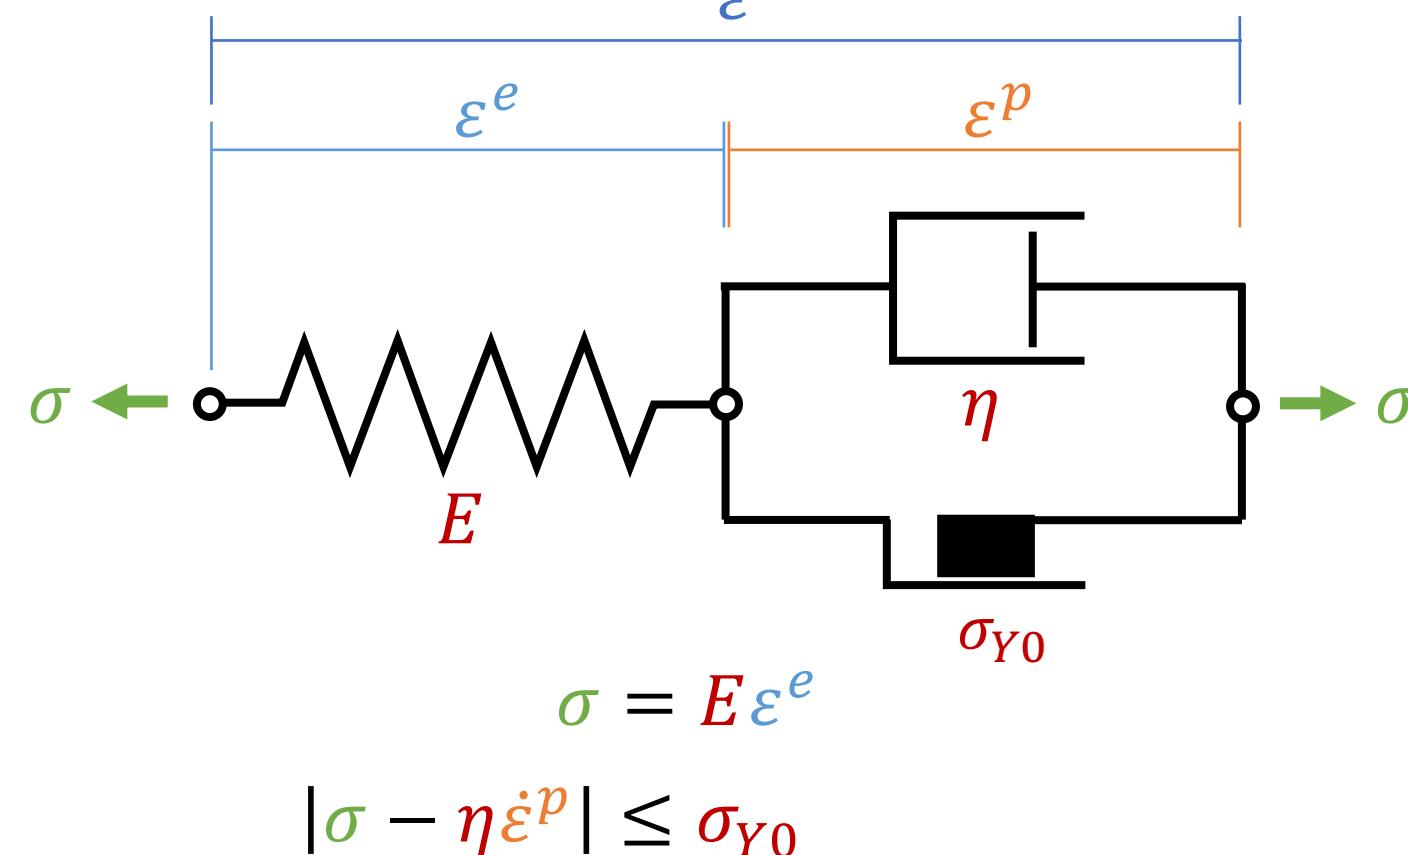
$$x_i = \hat{x}_i - \sum_{j > i} \textcolor{blue}{U}_{ij} x_j$$

$$\mathbf{x} = \textcolor{blue}{U}^{-1} \hat{\mathbf{x}}$$

$$\mathbf{x} = A^{-1} \mathbf{y}$$

The proposed set of bit-reversible operations can be used to implement reversible time-integrators for common visco-elastic/plastic constitutive models

Bingham-Maxwell Model:



$$\lambda = e^{-E\Delta t/\eta}$$

$$\varepsilon_0^p = \varepsilon - \text{sign}(\varepsilon - \varepsilon^p)\sigma_{Y0}/E$$

Forward:

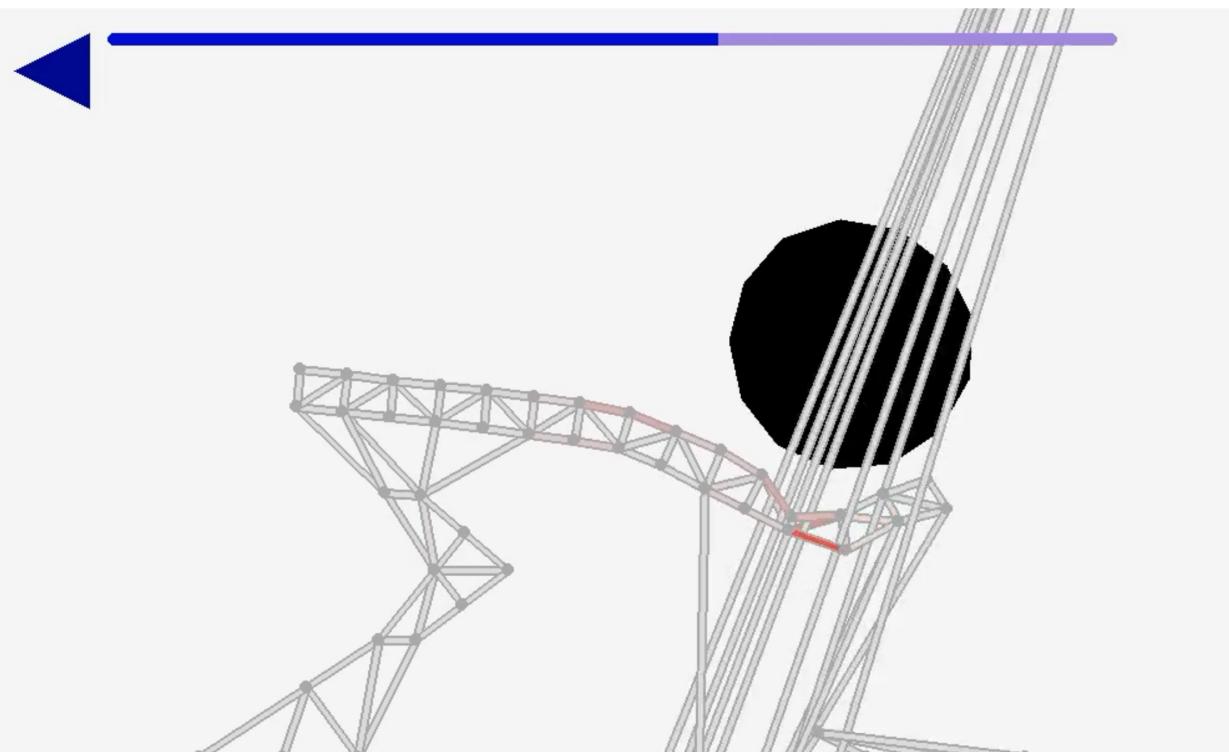
If $|E(\varepsilon - \varepsilon^p)| > \sigma_{Y0}$
 $(\varepsilon^p, \varepsilon^{p*}) \leftarrow (\varepsilon^p, \varepsilon^{p*}) [\times, \div] \lambda$
 $\varepsilon^p \leftarrow \varepsilon^p + \varepsilon_0^p(1 - \lambda)$

Backward:

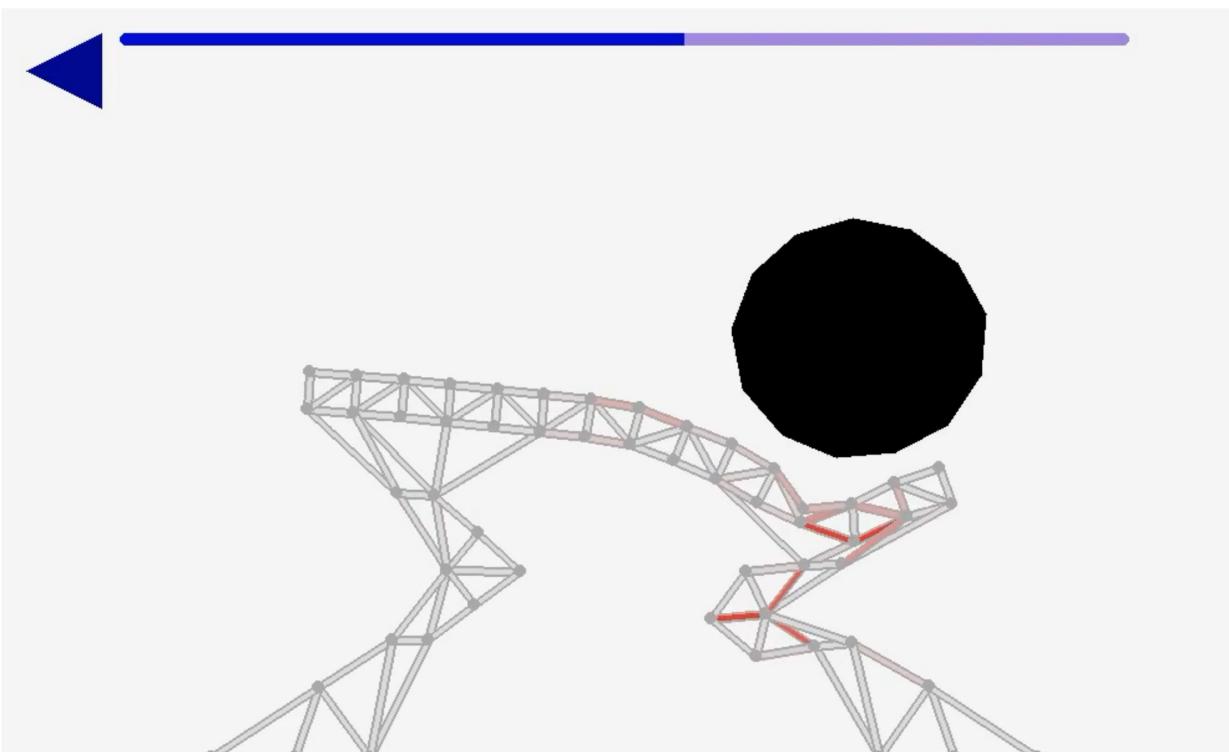
If $|E(\varepsilon - \varepsilon^p)| > \sigma_{Y0}$
 $\varepsilon^p \leftarrow \varepsilon^p - \varepsilon_0^p(1 - \lambda)$
 $(\varepsilon^p, \varepsilon^{p*}) \leftarrow (\varepsilon^p, \varepsilon^{p*}) [\div, \times] \lambda$

Demonstration: implementation of a reversible uniaxial visco-plasticity model with a maximum plastic strain-based failure criterion

Floating-point (*not* reversible)



Fixed-point (reversible)



Ongoing and future work

- Limitations:
 - Overflow!
 - Not all models are amenable to a reversible implementation
 - Must ensure *consistent* execution during forward/backward passes
- Alternative bit-roundoff data compression methods
- Inelastic material behavior
 - Continuum damage/plasticity/visco-elasticity, fracture, friction
- Compare performance on GPUs
 - Does the proposed approach help with minimizing device I/O and latency?
- Application to optimization/inverse problems
 - Optimal design of impact-resistant structures

Questions?